

① Shuffling (= uniformly random permutation)

Given a list of n elements, randomly permute
Specifically, want each possible permutation of
 $[n] = (1, 2, \dots, n)$ to have an equal chance, i.e.,
 $\frac{1}{n!}$ chance.

Naively: draw n iid samples from any absolutely
cts. probability distribution (eg. Gaussian)
then sort these.

Downside: a sort costs $O(n \cdot \log n)$ flops.

Can we get a linear time algo.?

Yes! "Knuth-shuffle" aka "Fisher Yates shuffle"

Input $X = (x_1, x_2, \dots, x_n)$

For $i = 1, 2, \dots, n-1$

$j \in \{i, i+1, \dots, n\}$ chosen uniformly at random

exchange x_i and x_j

time to compute this
is independent of i

Then: this is a uniform random permutation,

i.e., $P(\text{any permutation on } [n]) = \frac{1}{n!}$

Proof via induction, aka "loop invariant" in CS terms

def a " k -permutation" on $[n]$ is a list of k elements,
each element from $[n]$ and no repetitions.

i.e., the 1st k elements of any permutation

There are $\frac{n!}{(n-k)!}$ of these ($n!$ total permutations,
only count those whose
1st k elements differ)

So... we want a n -permutation.

claim: after iteration K , all K -permutations have an equal chance of equalling $\underbrace{X(1:K)}_{\text{Matlab notation}}$, i.e., $\frac{(n-K)!}{n!}$ chance, since there are $n!/(n-K)!$ K -perms.

proof-of-claim:

iter 1: there are n 1-permutations. By construction, all are equally likely.

iter $K+1$: (induction step)

Let $(X_{\sigma_1}, X_{\sigma_2}, \dots, X_{\sigma_K})$ be any $K+1$ permutation.

$$P(X(1:K+1) = (X_{\sigma_1}, \dots, X_{\sigma_K}, X_{\sigma_{K+1}}))$$

$$= P(X(K+1) = X_{\sigma_{K+1}} \mid X(1:K) = (X_{\sigma_1}, \dots, X_{\sigma_K})) \cdot$$

$$P(X(1:K) = (X_{\sigma_1}, \dots, X_{\sigma_K})) \rightarrow \underbrace{= \frac{(n-K)!}{n!}}_{\text{induct.}} \quad \swarrow$$

$X_{\sigma_{K+1}}$ is some entry not already selected

In algo, j is chosen uniformly at random among $\{K+1, \dots, n\}$ $n-K$ entries that have not already been selected.

So, this is $\frac{1}{n-K}$

$$\dots = \frac{1}{n-K} \cdot \frac{(n-K)!}{n!} = \frac{(n-(K+1))!}{n!}$$

==

So, after $K=n$ iterations, $\frac{0!}{n!} = \frac{1}{n!}$ chance of any given permutation,

i.e., this is a shuffle. \square

② Reservoir Sampling to form a Simple Random Sample (SRS) "randsample(n, k)" in Matlab

i.e. combination
not
permutation
not ordered
list

Def A "SRS" of k elements (from n possible) is
a subset of size k from $[n]$ such that all such
subsets have equal chance, namely $1/\binom{n}{k} = \frac{k!(n-k)!}{n!}$

"overkill"
since these
also
shuffle

Example: • shuffle the data and take 1st k entries
(inefficient)

- do the 1st k steps of Knuth / Fisher-Yates
- a SRS followed by a k -element shuffle is a k -perm.

Vitter '85
but
also '60's

Reservoir Sampling is for the situation where

1) we want only 1-pass over the data,
eg a data stream, aka streaming.

- 2a) n is unknown! ← classical
- or 2b) other things unknown! ← applicable to our class

Start here:

2a) First, observe Fisher-Yates looped $i=1, 2, \dots, n-1$
 $j \in \{i, \dots, n\}$

but we can do an "inside-out" version

- Input $x \in \mathbb{R}^n$, output $y \in \mathbb{R}^n$ (shuffled), no longer in-place
- For $i=1, 2, \dots, n-1$, $n \leftarrow$ different!

$j \in \{1, 2, \dots, i\}$ uniformly at random
(not $\{i, i+1, \dots, n\}$ anymore)

if $j \neq i$,

$y_i \leftarrow y_j$

$y_j \leftarrow x_i$

} if we initialized $y = x$,
then this is like
swap (y_i, y_j)
... and running outer loop
backward!

(we won't show it here, but similar inductive proof shows this is correct) i.e., reverse input, then flip "for" loop.

Observe that since $j \in \{1, \dots, i\}$ not $\{1, \dots, n\}$, we don't need to know n .

This is essentially the basic reservoir sampling algo, "Algo R" (Alan Waterman)

Input: $x \in \mathbb{R}^n$, n "unknown"

Output: $y \in \mathbb{R}^k$, y is the "reservoir"

initialize $y = x(1:k)$

For $i = k+1, \dots, n$ ← only place n appears

$j \in \{1, 2, \dots, i\}$ uniformly at random

if $j \leq k$

$y_j = x_i$

} if $j > k$, we'd be updating y past its last entry. We don't care, so skip it.

- one-pass ✓
- independent of n ✓
- Running time: $O(n)$ OK, though "Algo L" is $O(k + k \log(n/k))$ w/ tricks.
- k -perm? No, not shuffled (easy to see if $n=k$)
- SRS? yes.

2b) Let's keep k of n items (order still unimportant) such that x_i is kept with probability proportional to its weight, $w_i = w(x_i) \geq 0$,
 i.e., $w_i / W = \frac{n}{\sum_{j=1}^n w_j}$.

New complication: if n is unknown, so is W .

Ex: Sample a row of a matrix $X(i, :)$ proportional to its l_2 -norm squared. Then $W = \|X\|_F^2$.
 If we stream rows, W keeps increasing.

Note: not all weights realizable,

ex. $k=2, n=3$, wts = $[1, 0, 0]$

or $k=n$, must have

wts = uniform

incompatible w, $k=2$

→ link on wikipedia

Algo A-Chao (cf. Wikipedia, or M. Chao '82)
 or P. Efrimidis '15

Input: $x \in \mathbb{R}^n$

Output: $y \in \mathbb{R}^k$

initialize $y = x(1:k)$

$$W = \sum_{i=1}^k w(y_i)$$

For $i = k+1, k+2, \dots, n$

$$W \leftarrow W + w(x_i)$$

$$p = \frac{k \cdot w(x_i)}{W}$$

with prob. p , keep this sample x_i by assigning it to

y_j with $j \in [k]$ chosen uniformly at random

else, do nothing.

counter-intuitive.

Not obvious!

③ Different types of sampling (about) K elements from $[n]$

① SRS (w/o replacement): choose subset $\Omega \subseteq [n]$ of size K such that all subsets equally likely: "uniform"

let Ω_K be such a set

pros: nice conceptually, get right size, no duplicates

cons: no longer independent (but exchangeable)

uniform in diff. ways

② SRS w/ replacement:

choose a list $y \in \mathbb{R}^K$ s.t. each $y_i \sim \text{Unif}([n])$ i.i.d.

pros: independent, get right size

cons: may contain duplicates which seems like a waste

③ "Bernoulli": keep each x_i w/ probability $K/n =: p$

Expect to keep K entries total

pros: independent, no duplicates

cons: the size of our sample is only K on average, not deterministically

Relations

We'll be looking at $P(\text{Failure}(\Omega))$ of some randomized algo. that uses samples Ω

Formalize:

Prop (3.1 in Recht "A Simplex Approach...")

$P(\text{failure via } ①) \leq P(\text{failure via } ③)$ if $P(\text{failure}(\Omega_K)) \geq$

Proof Let Ω' be sampled via ② w/ K entries.

$P(\text{failure}(\Omega'_K))$

i.e., $\Omega' = \text{unique}(y)$

whenever $K \leq K'$
(more is better)

↑
set of length $\leq K$
K list of length K

$$\begin{aligned}
 \text{then } P(\text{failure}(\Omega')) &= \sum_{i=0}^K \underbrace{P(\text{fail}(\Omega') \mid |\Omega'|=i)}_{P(\text{fail}(\Omega_i)) \text{ is main observation}} \cdot P(|\Omega'|=i) \\
 &\geq P(\text{fail}(\Omega_K)) \text{ since } i \leq K \\
 &\geq P(\text{fail}(\Omega_K)) \cdot \underbrace{\sum_{i=0}^K P(|\Omega'|=i)}_{=1} \\
 &\geq P(\text{fail}(\Omega_K)). \quad \square
 \end{aligned}$$

Prop (p.15 "Robust Uncertainty Principles..." Candès et al '05)

Let Ω_K be sampled via ① and let

Ω' be sampled via ③ w/ parameter $p = K/n$ so $\mathbb{E}|\Omega'| = K$

then if $P(\text{failure}(\Omega_K))$ is non-increasing in K (same as in prev. prop.),

$$P(\text{failure}(\Omega')) \leq 2 \cdot P(\text{failure}(\Omega_K)).$$

[see Appendix of "Robust Principal Component Analysis" Candès et al. '09 for more refined 2-way result]

proof

$$P(\text{failure}(\Omega')) = \sum_{i=0}^n \overset{\text{new!}}{P(\text{fail}(\Omega') \mid |\Omega'|=i)} \cdot P(|\Omega'|=i)$$

(same as before, but sum goes to n , not K)

$$\geq \sum_{i=0}^K P(\dots \mid \dots) \cdot P(\dots)$$

since non-increasing

$$\geq \underbrace{P(\text{fail}(\Omega_K))}_{\text{as before}} \cdot \underbrace{\sum_{i=0}^K P(|\Omega'|=i)}_{\leq 1 \text{ now.}}$$

$p \cdot n = K \in \mathbb{Z}$

is the median

of $|\Omega'|$ so $P(|\Omega'| \leq K-1) < 1/2 < P(|\Omega'| \leq K)$ cf. Japles + Samuels '68

$$\text{so } \sum_{i=0}^K P(|\Omega'|=i) = P(|\Omega'| \leq K) > 1/2$$

$$\text{so } \dots > 1/2 P(\text{fail}(\Omega_K)). \quad \square$$